

## Deployment Strategies – Theory & Step-by-Step Execution

Understanding deployment strategies helps ensure **high availability, low risk, and smooth releases** in production environments.

Below are the **4 most important deployment strategies**, explained with **theory + practical steps** 

### ◆ Rolling/Rollout Deployment

#### Theory:

Rolling deployment updates the application **incrementally** by replacing instances **one by one**, while keeping the service available.

Old and new versions run **simultaneously** during deployment.

#### △ Key Point:

- **Zero/minimal downtime**
- Default strategy in Kubernetes

#### Steps:

##### 1. Create a Deployment with 3 replicas with any older image

```
controlplane:~$ kubectl create deployment mydep --image=nginx:1.24 --replicas=3
deployment.apps/mydep created
controlplane:~$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
mydep-7894647bfd-dfbct  1/1     Running   0          15s
mydep-7894647bfd-gc6js  1/1     Running   0          15s
mydep-7894647bfd-sxd7k  1/1     Running   0          15s
controlplane:~$
```

##### 2. Now update or change the version of image in that Deployment

```
controlplane:~$ kubectl set image deployments mydep nginx=docker.io/nginx:1.25
deployment.apps/mydep image updated
```

##### 3. And all pods will terminate one by one and new pods will be created

- As shown in below image

```
controlplane:~$ kubectl get pods -w
NAME           READY   STATUS    RESTARTS   AGE
mydep-689c7c668b-q8zkq  0/1    ContainerCreating   0    6s
mydep-7894647bfd-dfbct 1/1    Running   0    73s
mydep-7894647bfd-gc6js  1/1    Running   0    73s
mydep-7894647bfd-sxd7k  1/1    Running   0    73s
mydep-689c7c668b-q8zkq  1/1    Running   0    8s
mydep-7894647bfd-sxd7k  1/1    Terminating   0    75s
mydep-7894647bfd-sxd7k  1/1    Terminating   0    75s
mydep-689c7c668b-zg4sm  0/1    Pending   0    0s
mydep-689c7c668b-zg4sm  0/1    Pending   0    0s
mydep-689c7c668b-zg4sm  0/1    ContainerCreating  0    0s
mydep-7894647bfd-sxd7k  1/1    Terminating   0    76s
mydep-689c7c668b-zg4sm  0/1    ContainerCreating  0    1s
mydep-7894647bfd-sxd7k  0/1    Completed   0    76s
mydep-689c7c668b-zg4sm  1/1    Running   0    2s
mydep-7894647bfd-sxd7k  0/1    Completed   0    77s
mydep-7894647bfd-sxd7k  0/1    Completed   0    77s
mydep-7894647bfd-dfbct  1/1    Terminating   0    77s
mydep-689c7c668b-kkzp9  0/1    Pending   0    0s
mydep-7894647bfd-dfbct  1/1    Terminating   0    77s
mydep-689c7c668b-kkzp9  0/1    Pending   0    0s
mydep-689c7c668b-kkzp9  0/1    ContainerCreating  0    1s
mydep-7894647bfd-dfbct  1/1    Terminating   0    78s
mydep-7894647bfd-dfbct  0/1    Completed   0    78s
mydep-689c7c668b-kkzp9  0/1    ContainerCreating  0    1s
mydep-7894647bfd-dfbct  0/1    Completed   0    78s
mydep-7894647bfd-dfbct  0/1    Completed   0    78s
mydep-689c7c668b-kkzp9  1/1    Running   0    2s
mydep-7894647bfd-gc6js  1/1    Terminating   0    79s
mydep-7894647bfd-gc6js  1/1    Terminating   0    79s
mydep-7894647bfd-gc6js  1/1    Terminating   0    80s
mydep-7894647bfd-gc6js  0/1    Completed   0    80s
mydep-7894647bfd-gc6js  0/1    Completed   0    80s
```

## ◆ Recreate Deployment

### Theory:

Recreate deployment replaces the **entire old application version** with the **new version** by stopping all running instances first.

During deployment, the application remains **unavailable**.

### Steps:

#### 1. Create a YAML file with following details

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: mydep
  name: mydep
spec:
  replicas: 3
  selector:
    matchLabels:
      app: mydep
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: mydep
    spec:
      containers:
        - image: docker.io/nginx:1.24
          name: nginx
          resources: {}
status: {}
```

- Add strategy type = Recreate in spec.

## 2. Apply the YAML file and create a Deployment

```
controlplane:~$ kubectl apply -f mydep.yaml
deployment.apps/mydep configured
controlplane:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
mydep-54f78b9b99-bvgm6   1/1     Running   0          6s
mydep-54f78b9b99-kbhsv   1/1     Running   0          6s
mydep-54f78b9b99-qd8rd   1/1     Running   0          101s
controlplane:~$
```

## 3. Now update the image with new version in Deployment

```
controlplane:~$ kubectl set image deployments mydep nginx=nginx:1.25
deployment.apps/mydep image updated
controlplane:~$
```

## 4. Check pods by “kubectl get pods -w”

- All pod will get terminated at once and New pod will get created

```
controlplane:~$ kubectl get pods -w
NAME          READY   STATUS    RESTARTS   AGE
mydep-7894647bfd-hzw9k   1/1     Running   0          15s
mydep-7894647bfd-t78xf   1/1     Running   0          15s
mydep-7894647bfd-zqfp4   1/1     Running   0          15s
mydep-7894647bfd-hzw9k   1/1     Terminating   0          23s
mydep-7894647bfd-zqfp4   1/1     Terminating   0          23s
mydep-7894647bfd-t78xf   1/1     Terminating   0          23s
mydep-7894647bfd-hzw9k   1/1     Terminating   0          23s
mydep-7894647bfd-zqfp4   1/1     Terminating   0          23s
mydep-7894647bfd-t78xf   1/1     Terminating   0          23s
mydep-7894647bfd-hzw9k   1/1     Terminating   0          23s
mydep-7894647bfd-hzw9k   0/1     Completed   0          23s
mydep-7894647bfd-t78xf   1/1     Terminating   0          23s
mydep-7894647bfd-zqfp4   1/1     Terminating   0          23s
mydep-7894647bfd-t78xf   0/1     Completed   0          23s
mydep-7894647bfd-zqfp4   0/1     Completed   0          23s
mydep-7894647bfd-zqfp4   0/1     Completed   0          24s
mydep-7894647bfd-zqfp4   0/1     Completed   0          24s
mydep-7894647bfd-hzw9k   0/1     Completed   0          24s
mydep-7894647bfd-hzw9k   0/1     Completed   0          24s
mydep-7894647bfd-t78xf   0/1     Completed   0          24s
mydep-7894647bfd-t78xf   0/1     Completed   0          24s
mydep-867d69ffc5-kpzsl  0/1     Pending    0          0s
mydep-867d69ffc5-kpzsl  0/1     Pending    0          0s
mydep-867d69ffc5-zh9js  0/1     Pending    0          0s
mydep-867d69ffc5-7wr8x  0/1     Pending    0          0s
mydep-867d69ffc5-zh9js  0/1     Pending    0          0s
mydep-867d69ffc5-kpzsl  0/1     ContainerCreating  0          0s
mydep-867d69ffc5-7wr8x  0/1     Pending    0          0s
mydep-867d69ffc5-zh9js  0/1     ContainerCreating  0          0s
mydep-867d69ffc5-7wr8x  0/1     ContainerCreating  0          0s
mydep-867d69ffc5-kpzsl  0/1     ContainerCreating  0          0s
mydep-867d69ffc5-zh9js  0/1     ContainerCreating  0          1s
mydep-867d69ffc5-7wr8x  0/1     ContainerCreating  0          1s
mydep-867d69ffc5-7wr8x  1/1     Running    0          2s
mydep-867d69ffc5-zh9js  1/1     Running    0          2s
mydep-867d69ffc5-kpzsl  1/1     Running    0          2s
```

## ⚠ Key Point:

- Causes **downtime**
- Simple but risky for production

## ◆ Canary Deployment

### Theory:

Canary deployment releases the new version to a **small group of users first**, then gradually increases traffic if everything works correctly.

### Steps:

#### 1. Create a YAML file with following details

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: mycan1
    name: mycan1
spec:
  replicas: 7
  selector:
    matchLabels:
      app: mycan1
      class: unnati
  strategy: {}
  template:
    metadata:
      labels:
        app: mycan1
        class: unnati
    spec:
      containers:
        - image: hashicorp/http-echo
          name: http-echo
          args:
            - "-text=version1"
            - "-listen=:5678"
          resources: {}
status: {}
```

- **Class = unnati**
- **args = version1**
- **Replicas = 7**

#### 2. Apply the YAML file and create a deployment “mycan1”

```

controlplane:~$ kubectl apply -f mycan1.yaml
deployment.apps/mycan1 unchanged
controlplane:~$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
mycan1-797f74f8d-494n6   1/1     Running   0          17s
mycan1-797f74f8d-h2zf5   1/1     Running   0          17s
mycan1-797f74f8d-jcwsx   1/1     Running   0          17s
mycan1-797f74f8d-szb8p   1/1     Running   0          17s
mycan1-797f74f8d-tmtxj   1/1     Running   0          17s
mycan1-797f74f8d-v5thv   1/1     Running   0          17s
mycan1-797f74f8d-xmzxt   1/1     Running   0          17s
controlplane:~$ █

```

**3. Create a 2nd YAML file with the same details only change in name and args use same class = unnati.**

```

apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: mycan2
    name: mycan2
spec:
  replicas: 3
  selector:
    matchLabels:
      app: mycan2
      class: unnati
  strategy: {}
  template:
    metadata:
      labels:
        app: mycan2
        class: unnati
    spec:
      containers:
        - image: hashicorp/http-echo
          name: http-echo
          args:
            - "-text=version2"
            - "-listen=:5678"
          resources: {}
  status: {}

```

- **Replicas = 2**
- **Args = version2**
- **Class = unnati**

```

controlplane:~$ kubectl apply -f mycan2.yaml
deployment.apps/mycan2 created
controlplane:~$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
mycan1-797f74f8d-494n6  1/1    Running   0          5m24s
mycan1-797f74f8d-h2zf5  1/1    Running   0          5m24s
mycan1-797f74f8d-jcwsx  1/1    Running   0          5m24s
mycan1-797f74f8d-szb8p  1/1    Running   0          5m24s
mycan1-797f74f8d-tmtxj  1/1    Running   0          5m24s
mycan1-797f74f8d-v5thv  1/1    Running   0          5m24s
mycan1-797f74f8d-xmzxt  1/1    Running   0          5m24s
mycan2-84c5d47fb8-4d9w2  1/1    Running   0          7s
mycan2-84c5d47fb8-t5vmz  1/1    Running   0          7s
mycan2-84c5d47fb8-v57bg  1/1    Running   0          7s
controlplane:~$ █

```

#### 4. Create SVC with following details.

```

apiVersion: v1
kind: Service
metadata:
  labels:
    class: unnati
    name: mycansvc
spec:
  ports:
  - port: 5678
    protocol: TCP
    targetPort: 5678
  selector:
    class: unnati
status:
  loadBalancer: {}

```

- Give label as **class = unnati**

```

controlplane:~$ kubectl apply -f svc.yaml
service/mycansvc created
controlplane:~$ █

```

**4. Copy the SVC ip and “curl” it and you can see that 70% traffic git hitting version1 (mycan1) and 30% to version2 (mycan2).**

## ⚠ Key Point:

- **Lowest risk strategy**
  - Needs monitoring & traffic control tools

## ◆ Blue-Green Deployment

### Theory:

Blue-Green deployment maintains **two identical environments**:

- **Blue** → current production version
- **Green** → new application version

Traffic is switched to Green **only after successful testing**.

### Steps:

#### 1. Create Blue deployment with following details

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: blue
spec:
  replicas: 2
  selector:
    matchLabels:
      app: httpd
      version: blue
  template:
    metadata:
      labels:
        app: httpd
        version: blue
    spec:
      containers:
        - name: httpd
          image: httpd:2.4
          ports:
            - containerPort: 80
          volumeMounts:
            - name: html-volume
              mountPath: /usr/local/apache2/htdocs/index.html
              subPath: index.html
      volumes:
        - name: html-volume
      configMap:
        name: blue-html
```

#### 2. Create Green deployment with following details.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: green
spec:
  replicas: 2
  selector:
    matchLabels:
      app: httpd
      version: green
  template:
    metadata:
      labels:
        app: httpd
        version: green
    spec:
      containers:
        - name: httpd
          image: httpd:2.4
          ports:
            - containerPort: 80
          volumeMounts:
            - name: html-volume
              mountPath: /usr/local/apache2/htdocs/index.html
              subPath: index.html
      volumes:
        - name: html-volume
      configMap:
        name: green-html
```

### 3. Create config map for blue deployment with following details.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: blue-html
data:
  index.html: |
    <html>
      <body>
        <h1 style="color:blue">Welcome to BLUE Deployment!</h1>
      </body>
    </html>
```

### 4. Create config map for green deployment with following details.

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: green-html
data:
  index.html: |
    <html>
      <body>
        <h1 style="color:green">Welcome to GREEN Deployment!</h1>
      </body>
    </html>

```

## 5. Get SVC ip and do “curl” you will able to see the Blue deployment is working

```

controlplane:~$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
httpd-svc  LoadBalancer  10.97.2.2      <pending>      80:31119/TCP  6m3s
kubernetes  ClusterIP   10.96.0.1      <none>        443/TCP      20d
mycansvc   ClusterIP   10.105.113.65  <none>        5678/TCP      34m
controlplane:~$ curl 10.97.2.2
<html>
  <body>
    <h1 style="color:blue">Welcome to BLUE Deployment!</h1>
  </body>
</html>
controlplane:~$ █

```

## 6. Now do changes in SVC change the selector to “green”

```

spec:
  allocateLoadBalancerNodePorts: true
  clusterIP: 10.97.2.2
  clusterIPs:
  - 10.97.2.2
  externalTrafficPolicy: Cluster
  internalTrafficPolicy: Cluster
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - nodePort: 31119
    port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: httpd
    version: green █
  sessionAffinity: None
  type: LoadBalancer
  status:
    loadBalancer: {}

```

## 7. Now you “curl” the same sec ip and you will able to see Green deployment.

```
controlplane:~$ kubectl edit svc httpd-svc
service/httpd-svc edited
controlplane:~$ curl 10.97.2.2
<html>
  <body>
    <h1 style="color:green">Welcome to GREEN Deployment!</h1>
  </body>
</html>
controlplane:~$
```

### ⚠ Key Point:

- Instant rollback
- Requires double infrastructure

### 🎯 Conclusion:

Each deployment strategy has its own **use case**:

- Recreate → Simplicity
- Rolling → Continuous availability
- Blue-Green → Safe & fast rollback
- Canary → Controlled risk

A good DevOps engineer chooses the **right strategy at the right time**.

